

# Linear Algebra: Part II

Math Tools: Week 2

# Overview

## 1. Eigenvectors and eigenvalues

1. Quick review of matrix-vector multiplication
2. General applications

## 2. PCA

1. Mathematical derivation
2. Example

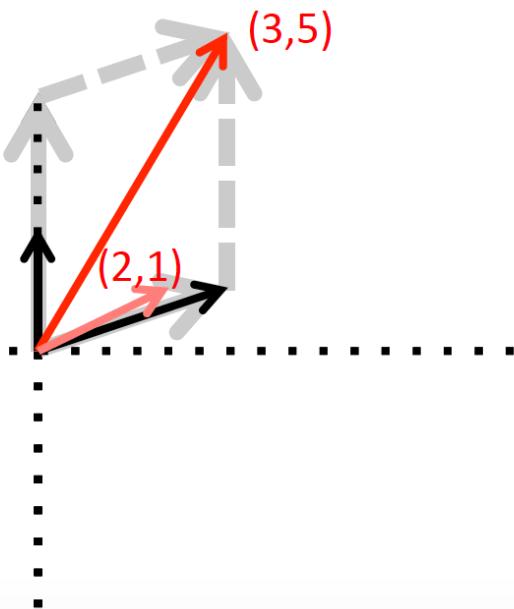
## 3. SVD

1. Applications

# Recall from last week: What do matrices do to vectors?

Outer product  
interpretation

$$\begin{bmatrix} 0 \\ 2 \\ 1 \end{bmatrix} \begin{bmatrix} 2 \\ 1 \end{bmatrix} = 2 \begin{bmatrix} 0 \\ 2 \end{bmatrix} + 1 \begin{bmatrix} 3 \\ 1 \end{bmatrix} = \begin{bmatrix} 0 \\ 4 \end{bmatrix} + \begin{bmatrix} 3 \\ 1 \end{bmatrix} = \begin{bmatrix} 3 \\ 5 \end{bmatrix}$$



Columns of M ‘span the plane’

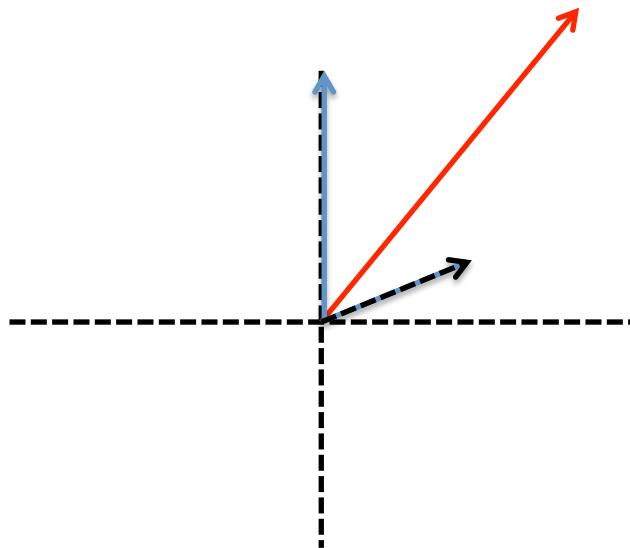
→ Different combinations of the columns of M can give you any vector in this plane

New vector is rotated and scaled

# Recall from last week: What do matrices do to vectors?

Inner product interpretation

$$\begin{bmatrix} 0 \\ 2 \\ 1 \end{bmatrix} \begin{bmatrix} 2 \\ 1 \end{bmatrix} = \begin{bmatrix} 0 & 3 \\ 2 & 1 \end{bmatrix} \begin{bmatrix} 2 \\ 1 \end{bmatrix} = \begin{bmatrix} 3 \\ 5 \end{bmatrix}$$



Rows of M also ‘span the plane’

Inner product multiplication can be thought of as a shorthand way to do several inner products

New vector is rotated and scaled

# Relating this to eigenvalues and eigenvectors

$$\overleftrightarrow{W} \vec{v} = \vec{u}$$

The matrix  $W$  transforms  $v$  into  $u$  – maps  $v$  onto  $u$ . Resulting vector can be a rotated and scaled version of  $v$ .

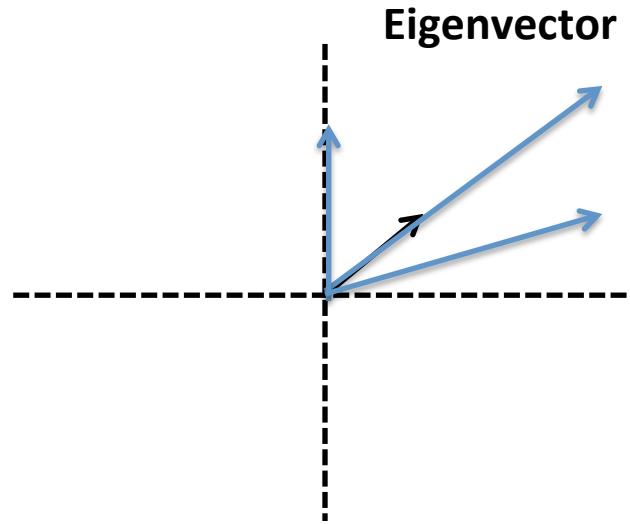
But it doesn't have to be rotated...

The special vectors for which the direction is preserved but the length changes:

## EIGENVECTOR

An example:

$$\begin{bmatrix} 0 & 3 \\ 2 & 1 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \end{bmatrix} = 1 \begin{bmatrix} 0 \\ 2 \end{bmatrix} + 1 \begin{bmatrix} 3 \\ 1 \end{bmatrix}$$
$$= \begin{bmatrix} 3 \\ 3 \end{bmatrix} = 3 \begin{bmatrix} 1 \\ 1 \end{bmatrix}$$

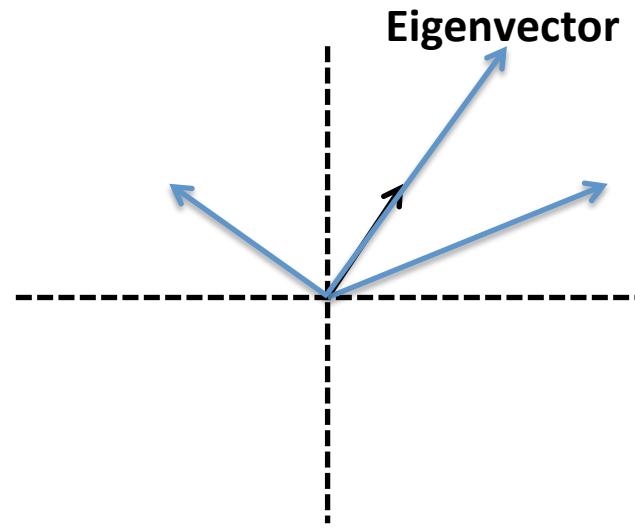


# Eigenvalues and eigenvectors

$$\begin{bmatrix} 0 & 3 \\ 2 & 1 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \end{bmatrix} = 3 \begin{bmatrix} 1 \\ 1 \end{bmatrix}$$

$$W\overrightarrow{v} = \lambda \overrightarrow{v}$$

**Eigenvector**      **Eigenvalue**



For an eigenvector, multiplication by the matrix is the same as multiplying the vector by a scalar. This scalar is called an eigenvalue.

How many eigenvectors can a matrix have???

- An  $n \times n$  matrix (the only type we will consider here) can have up to  $n$  distinct eigenvalues and  $n$  eigenvectors
- The set of eigenvectors, with distinct eigenvalues, are linearly independent. For symmetric matrices: they are orthogonal (their dot product is 0).

# Illustration of eigenvector and eigenvalues



Population dynamics of rabbits



$$x = \begin{bmatrix} \text{\# of rabbits b/w 0 and 1 year old} \\ \text{\# of rabbits b/w 1 and 2 years old} \\ \text{\# of rabbits b/w 2 and 3 years old} \end{bmatrix} \begin{bmatrix} 10 \\ 10 \\ 10 \end{bmatrix}$$

A = age transition matrix – probability that a member of the  $i^{\text{th}}$  age class will become a member of the  $(i+1)^{\text{th}}$  age class

$$\begin{bmatrix} 0 & 6 & 8 \\ 0.5 & 0 & 0 \\ 0 & 0.5 & 0 \end{bmatrix}$$

Population growth/loss:  $Ax_t = x_{t+1}$

$$\begin{bmatrix} 0 & 6 & 8 \\ 0.5 & 0 & 0 \\ 0 & 0.5 & 0 \end{bmatrix} \begin{bmatrix} 10 \\ 10 \\ 10 \end{bmatrix} = \begin{bmatrix} 140 \\ 5 \\ 5 \end{bmatrix}$$

In this case, an eigenvector of matrix A represents a ‘stable’ population distribution

$$Ax_t = \lambda x_{t+1} \quad x_t = x_{t+1}$$

If  $\lambda = 1$ , then there is a solution for which the population doesn’t change every year.

If  $\lambda < 1$ , then the population is shrinking

If  $\lambda > 1$ , then the population is growing

If the population starts on an eigenvector, it will stay on the eigenvector

This will be revisited in a few classes!

# Side-note on eigenvectors

Previous example

$$\overleftarrow{\overrightarrow{W}} \overrightarrow{v} = \lambda \overrightarrow{v}$$

discusses RIGHT eigenvectors – the eigenvector is on the RIGHT side of the matrix.

Does this matter? Yes, maybe. If a matrix is symmetric, the right and left eigenvectors are the same.

Left eigenvectors can be found by solving this equation:

$$\overrightarrow{v}^T \overleftarrow{\overrightarrow{W}} = \lambda \overrightarrow{v}$$

Note: This eigenvector is a row vector

Most of the time, left eigenvectors aren't that useful. But it's good to know that they exist.

# How to find eigenvalues and eigenvectors

Real life and easiest way: Using MATLAB's `eig` command

`[V,D] = eig(W)` will return a matrix  $V$ , where each column is an eigenvector, and a diagonal matrix  $D$  with the corresponding eigenvalues along the diagonal

MATLAB output:

`[V,D] = eig(W)`

$$W = \begin{bmatrix} 0 & 3 \\ 2 & 1 \end{bmatrix}$$

$V =$

-0.8321 -0.7071  
0.5547 -0.7071

Eigenvectors are defined only up to a scale factor, and so they are typically scaled such that the norm of the vector is 1

$D =$

-2 0  
0 3

Eigenvalue is 3, like we found before!

The set of eigenvalues is called the *spectrum* of a matrix

# How to find eigenvalues and eigenvectors

The math way:

Eigenvector-eigenvalue equation:  $\overleftrightarrow{W}\vec{v} = \lambda\vec{v}$

$$\overleftrightarrow{W}\vec{v} - \lambda\vec{v} = 0$$

$$(\overleftrightarrow{W} - \lambda I)\vec{v} = 0$$

(WHITEBOARD) Assume that  $v$  is not 0, this equation only has a solution if:

$$\det(\overleftrightarrow{W} - \lambda I) = 0$$

If  $\overleftrightarrow{W} = \begin{bmatrix} a & b \\ c & d \end{bmatrix}$  then  $\overleftrightarrow{W} - \lambda I = \begin{bmatrix} a - \lambda & b \\ c & d - \lambda \end{bmatrix}$

$$\begin{aligned} \det(\overleftrightarrow{W} - \lambda I) &= (a - \lambda)(d - \lambda) - bc \\ &= \lambda^2 - \lambda(a + d) + ad - bc \\ &= \lambda^2 - \lambda \text{Trace} + \text{Det} \end{aligned}$$

$$\lambda = \frac{\text{Trace} \pm \sqrt{\text{Trace}^2 - 4\text{Det}}}{2}$$

# Practice

Find eigenvectors and eigenvalues of this matrix:

$$W = \begin{bmatrix} 4 & -1 \\ 2 & 1 \end{bmatrix}$$

# Some fun facts that actually come in handy

$$W = \begin{bmatrix} a & b \\ c & d \end{bmatrix}$$

Determinant of a matrix is equal to the product of eigenvalues

$$ad - bd = \lambda_1 \lambda_2$$

Trace of a matrix is equal to the sum of eigenvalues

$$a + d = \lambda_1 + \lambda_2$$

Eigenvalues of a triangular matrix can be read off the diagonal:

$$\begin{bmatrix} 3 & 1 \\ 0 & 2 \end{bmatrix} \quad \lambda_1 = 3 \quad \lambda_2 = 2$$

# When are eigenvectors/eigenvalues useful?

1. **Gives you a way to ‘decompose a matrix’**
2. Allows some easy shortcuts in computation
3. Give you a sense of what kind of ‘matrix’ or dynamics you are dealing with
4. Allows for a convenient change of basis
5. Frequently used in both modeling and data analysis

# Eigen-decomposition (matrix diagonalization)

Eigen-decomposition, or the decomposition of a matrix into eigenvalues and eigenvectors, can be thought of as similar to prime factorization

Example with primes:  $12 = 2 \cdot 2 \cdot 3$

Example with matrices:  $W = VDV^{-1}$

$V$  is a matrix of eigenvectors, and  $D$  is a diagonal matrix of eigenvalues (WB)

Note: only works for NxN matrices with N linearly independent eigenvectors

Eigenvectors of a matrix form an orthogonal basis for the matrix – any other vector in the column space of the matrix can be re-written with eigenvectors:

General example:

$$\vec{v} = c_1 \vec{v}_1 + c_2 \vec{v}_2$$

$$\vec{u} = W \vec{v}$$

$$\vec{u} = W(c_1 \vec{v}_1 + c_2 \vec{v}_2)$$

$$\vec{u} = c_1 W \vec{v}_1 + c_2 W \vec{v}_2$$

$$W \vec{v}_1 = \lambda_1 \vec{v}_1$$

$$\vec{u} = c_1 \lambda_1 \vec{v}_1 + c_2 \lambda_2 \vec{v}_2$$

Don't even need the matrix!

# When are eigenvectors/eigenvalues useful?

Eigenvalues can reveal the preferred inputs to a system

$$\vec{u} = W \vec{v} \quad \vec{v} = c_1 \vec{v}_1 + c_2 \vec{v}_2 \quad \vec{u} = c_1 \lambda_1 \vec{v}_1 + c_2 \lambda_2 \vec{v}_2$$

Assume that all  $v$ 's are of unit length, so the length of  $u$  depends on the  $c$ 's and  $\lambda$ 's

If a vector  $v$  points in the direction of eigenvector  $v_1$ , then  $c_1$  will be large (or at least positive)

If  $\lambda_1$  is also large, then multiplication by  $W$  effectively lengthens the vector along this direction

More generally – the eigenvectors with the large eigenvalues can tell you which input vectors ( $v$ ) the matrix “prefers”, or will give a large response to

# Connect this back to Neuroscience

$$\vec{u} = W \vec{v}$$

$$\vec{u} = c_1 \lambda_1 \vec{v}_1 + c_2 \lambda_2 \vec{v}_2$$

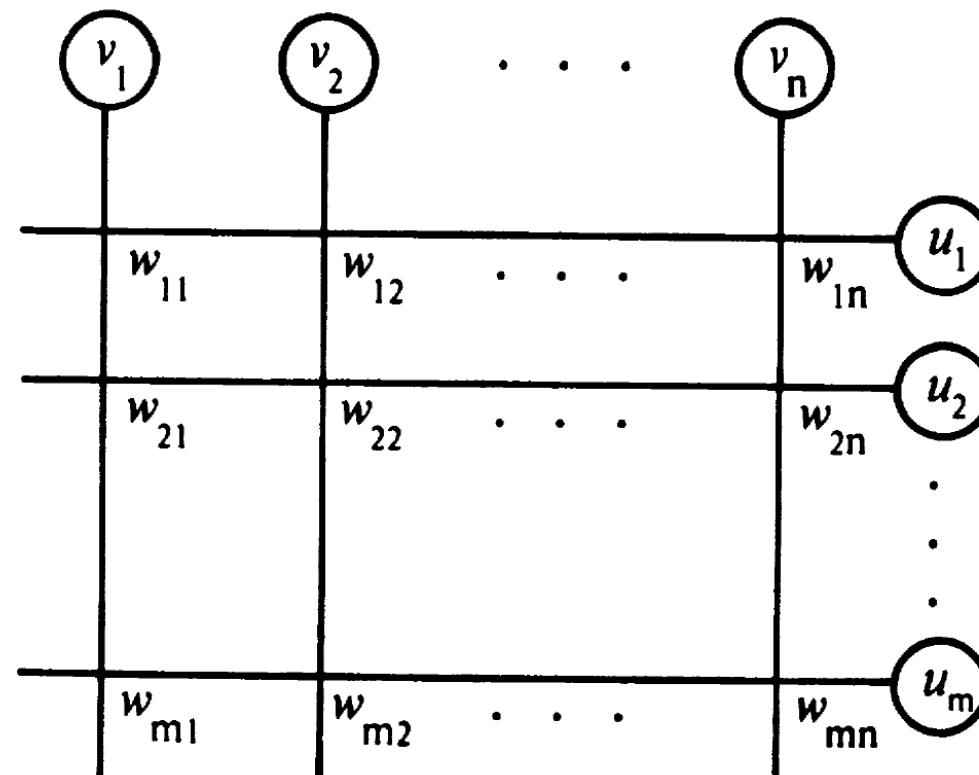
Each output value  $u$  is computed by an inner product with the input vector  $v$  and the weight matrix row,  $w$

Large activation happens when the input vector closely matches eigenvectors with large eigenvalues

Also: Given  $v$  and  $u$ , you can compute the matrix  $W$  (well-defined if  $u$  is long enough)

**Inputs**

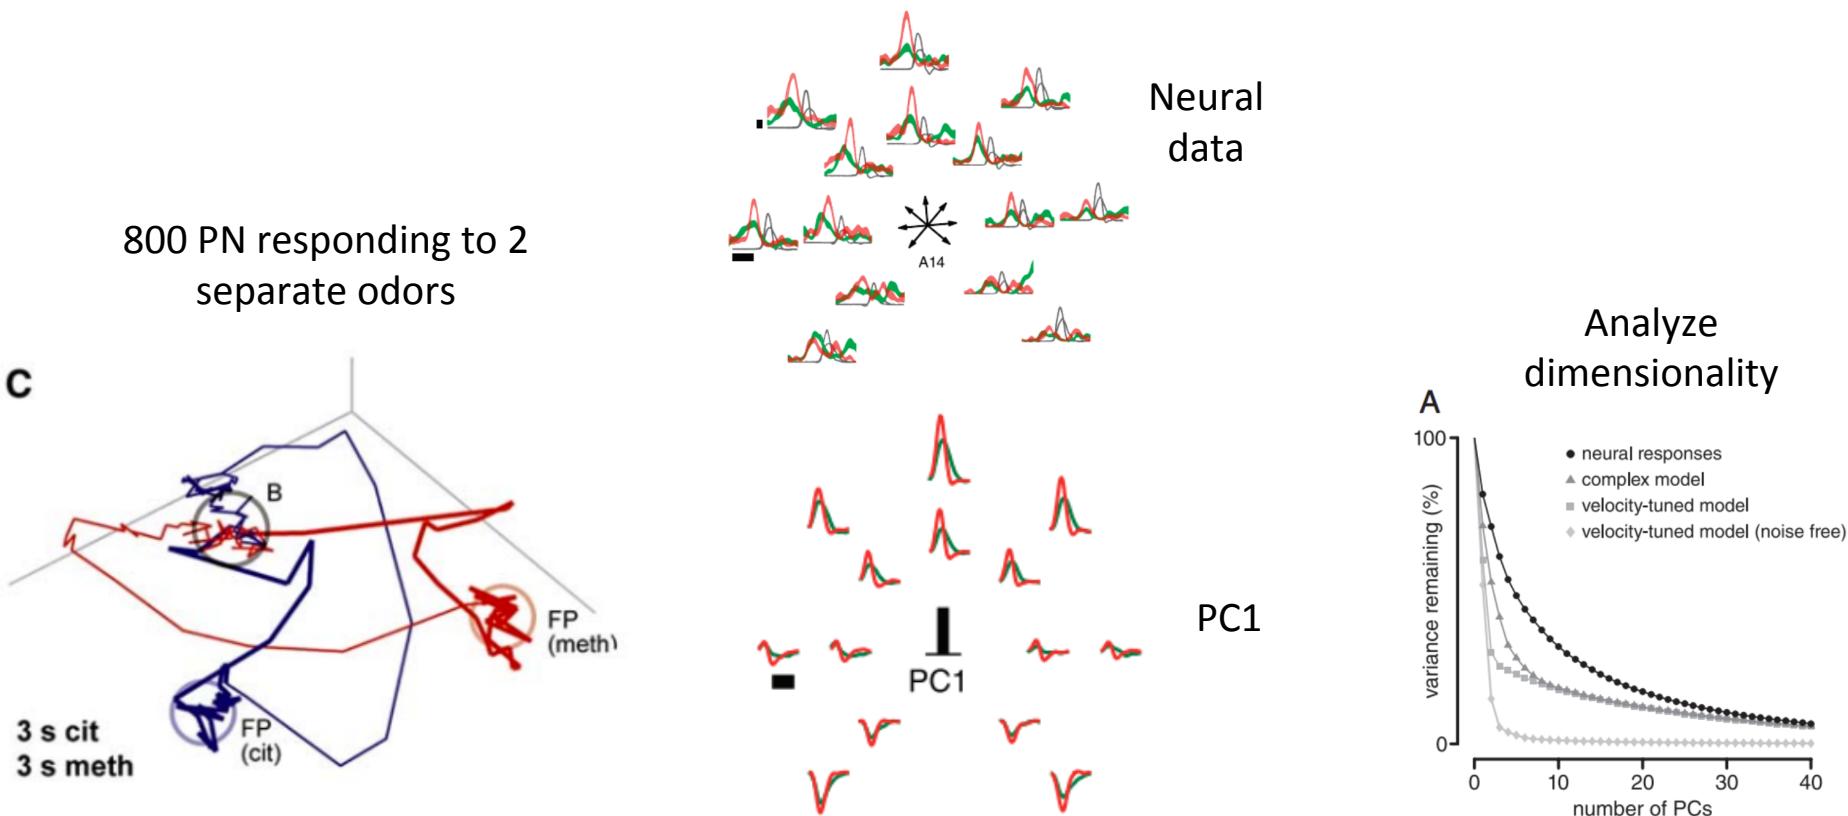
**Outputs**



# Utility of eigenvectors/eigenvalues: PCA

## What is PCA?

- Dimensionality reduction method (that keeps as much variation as possible)
- Pull structure out of seemingly complicated datasets
- De-noising technique
- Useful when # of variables that you record may be 'excessive'

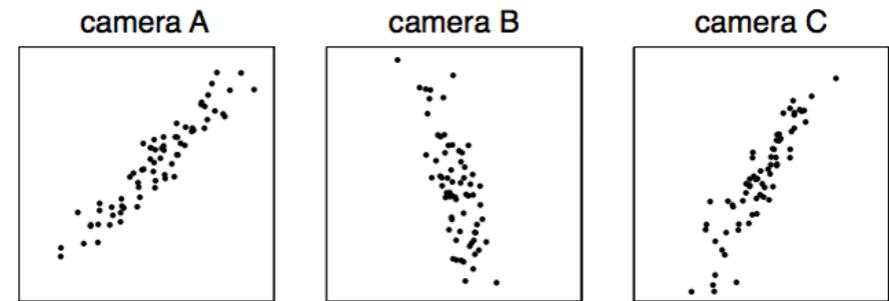
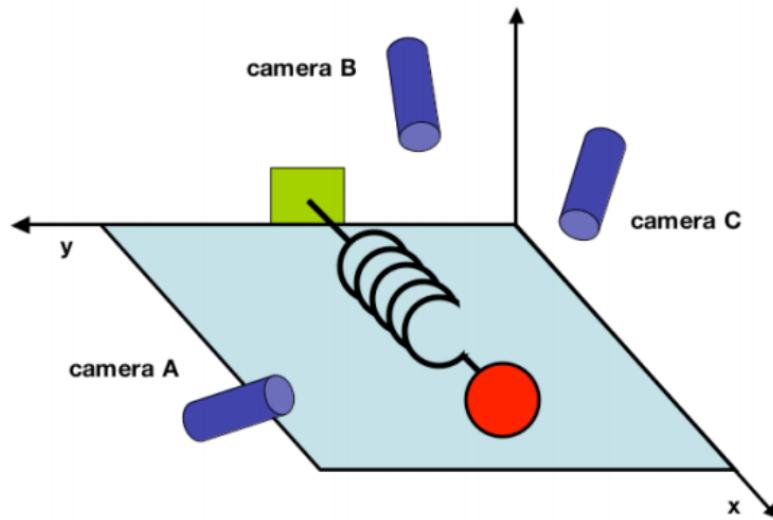


# PCA: a motivating example

We do an experiment: we record a spring (with a mass attached at the end) oscillate in one dimension:

BUT, we don't know this – don't know which measurements best reflect the dynamics of our system

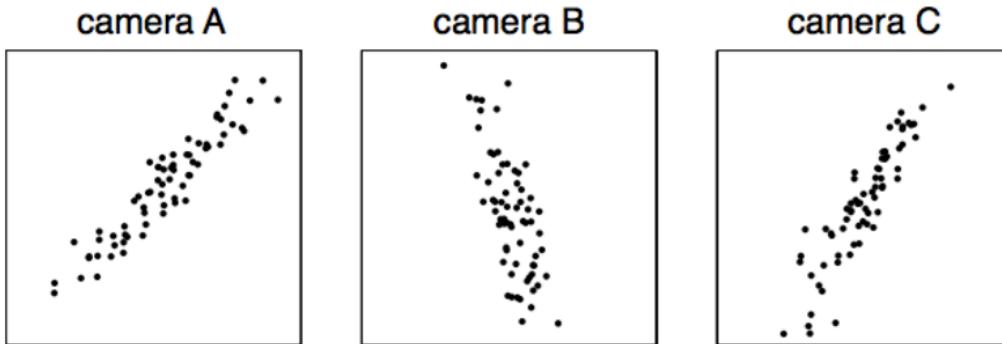
Record in 3 dimensions



Each black dot – position of ball at each time frame

Can we determine the dimension along which the dynamics occur? Want to uncover that the dynamics are truly 1-dimensional

# PCA: a motivating example



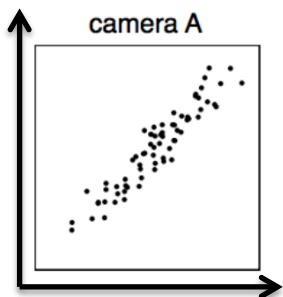
Data vector  
for 1 time  
point:

$$\vec{X} = \begin{bmatrix} x_A \\ y_A \\ x_B \\ y_B \\ x_C \\ y_C \end{bmatrix}$$

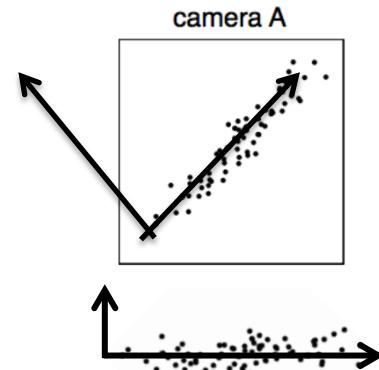
Together, data points form some cloud in a 6-dimensional space

6-d space has too many dimensions to visualize, and the data might not be truly 6-dimensional anyway

Is there a better way to represent this data?



Naïve basis for each  
camera:  $\{(0,1), (1,0)\}$



Could choose a different  
basis that makes more  
sense for this data

**Assumption/limit of PCA:** new basis will be a linear combination of the old basis

# PCA: into the details

Define the relevant matrices:

Data matrix:  $\mathbf{X} = \begin{bmatrix} \overrightarrow{X_1} & \overrightarrow{X_2} & \dots & \overrightarrow{X_T} \end{bmatrix}$

Each vector:  $\overrightarrow{X} = \begin{bmatrix} x_A \\ y_A \\ x_B \\ y_B \\ x_C \\ y_C \end{bmatrix}$

If we record for 10 mins at 120 Hz, X has 6 rows and  
 $10*60*120 = 72000$  columns

To project the data into a more sensible basis:

$$\mathbf{P} \mathbf{X} = \mathbf{Y}$$

↑  
Projection matrix      Data matrix      New data matrix  
6 x 6                  6 x 72000      6 x 72000

$\mathbf{P}$  transforms  $\mathbf{X}$  into  $\mathbf{Y}$  through a rotation and scaling

The *rows* of  $\mathbf{P}$  are the basis vectors for expressing the *columns* of  $\mathbf{X}$  – they are the principal components!

$$\mathbf{P} \mathbf{X} = \begin{bmatrix} \overrightarrow{p_1} \\ \overrightarrow{p_2} \\ \vdots \\ \overrightarrow{p_m} \end{bmatrix} \begin{bmatrix} \overrightarrow{X_1} & \overrightarrow{X_2} & \dots & \overrightarrow{X_T} \end{bmatrix} \quad \mathbf{Y} = \begin{bmatrix} \overrightarrow{p_1} \cdot \overrightarrow{X_1} & \dots & \overrightarrow{p_1} \cdot \overrightarrow{X_T} \\ \vdots & \ddots & \vdots \\ \overrightarrow{p_m} \cdot \overrightarrow{X_1} & \dots & \overrightarrow{p_m} \cdot \overrightarrow{X_T} \end{bmatrix}$$

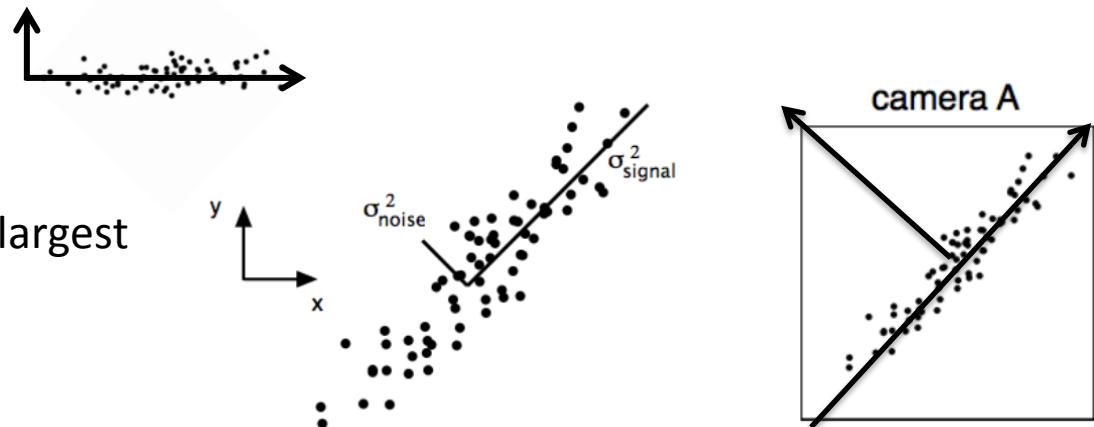
# PCA: into the details

$$\mathbf{P}\mathbf{X} = \mathbf{Y}$$

$$\mathbf{Y} = \begin{bmatrix} \vec{p_1} \cdot \vec{X}_1 & \dots & \vec{p_1} \cdot \vec{X}_T \\ \vdots & \ddots & \vdots \\ \vec{p_m} \cdot \vec{X}_1 & \dots & \vec{p_m} \cdot \vec{X}_T \end{bmatrix}$$

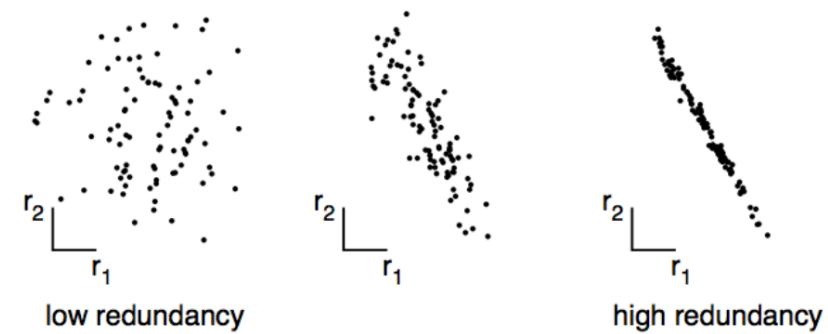
Okay... so how do we choose P?  
What do we want?

Assume that the direction with the largest variance comes from real dynamics



Look for directions that capture redundancy – might allow us to reduce the dimensions

Want  $\mathbf{Y}$  to have low redundancy



# PCA: into the details

Can quantify these goals through the covariance matrix!

For two vectors:

Variance:  $\sigma_A^2 = \frac{1}{n} \sum_i a_i^2$      $\sigma_B^2 = \frac{1}{n} \sum_i b_i^2$     (assume the data is mean-subtracted)

Covariance:  $\sigma_{AB}^2 = \frac{1}{n} \sum_i a_i b_i$     Rewriting using vector notation:

$$a = [a_1 \quad a_2 \quad \cdots \quad a_n]$$

$$b = [b_1 \quad b_2 \quad \cdots \quad b_n]$$

$$\sigma_{AB}^2 = \frac{1}{n} ab^T$$

For lotsa vectors:

$$Z = \begin{bmatrix} \vec{z}_1 \\ \vec{z}_2 \\ \vdots \\ \vec{z}_n \end{bmatrix}$$

Each row corresponds to measurements of a certain type (from one camera)

$$C_Z = \frac{1}{n} ZZ^T$$

Diagonal elements correspond to variance and off-diagonal elements correspond to covariance between measurement types

# PCA: into the details

$$C_Z = \frac{1}{n} ZZ^T$$

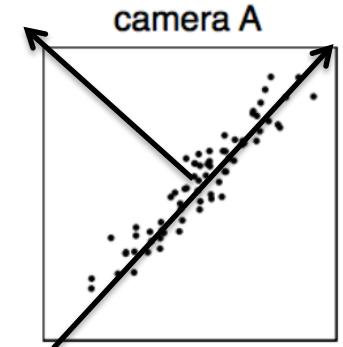
Diagonal elements correspond to variance within measurement types and off-diagonal elements correspond to covariance between measurement types

Want the covariance matrix for our projected data ( $\mathbf{Y}$ ) to have large values along the diagonal (large variance for the signal) and small values everywhere else (low covariance – low redundancy)

Technical term for this: diagonalize  $\mathbf{C}_Y$

One algorithm that accomplishes this:

1. Select a normalized direction in  $m$ -dimensional space along which the variance is maximized. This is  $p_1$ .
2. Find an orthogonal direction along which remaining variance is maximized. This is  $p_2$ .
3. Repeat until  $m$  vectors are chosen.



# PCA: eigenvector-based solution

Goal: Find some orthonormal matrix  $P$  for  $Y = PX$  such that  $C_Y = (1/n)YY^T$  is a diagonal matrix

$$\begin{aligned} C_Y &= \frac{1}{n}YY^T \\ &= \frac{1}{n}(PX)(PX)^T \\ &= \frac{1}{n}PXX^TP^T \quad C_x = \frac{1}{n}XX^T \\ &= PC_XP^T \end{aligned}$$

$C_x$  is symmetric, so we can decompose it into the following:  $C_X = EDE^T$

$E$  is a matrix of eigenvectors and  $D$  is a diagonal matrix with eigenvalues along the diagonal

$$C_Y = P(EDE^T)P^T$$

**What if we choose  $P = E^T$ ?**

$$C_Y = P(P^TDP)P^T \quad \text{Because } P \text{ is a matrix of orthogonal eigenvectors, } P^T = P^{-1}$$

$$C_Y = (PP^{-1})D(PP^{-1})$$

$$C_Y = D$$

Our choice of  $P$  diagonalized  $C_Y$ !

# PCA: summary and utility

What just happened:

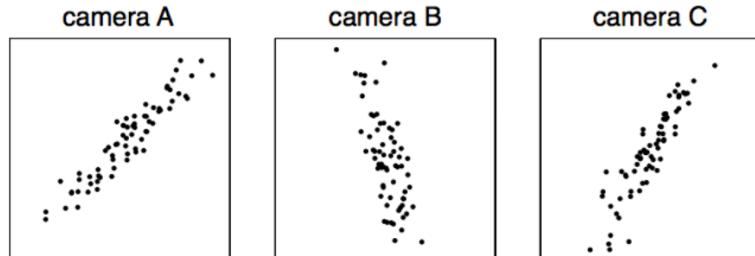
If we choose  $P$  to be the matrix of eigenvectors of the covariance matrix of  $X$ , then  $Y = PX$  returns a  $Y$  matrix whose covariance matrix is diagonalized.

- The principal components (or ‘axes’) are the eigenvectors of the covariance matrix of the data matrix  $X$

Now what???

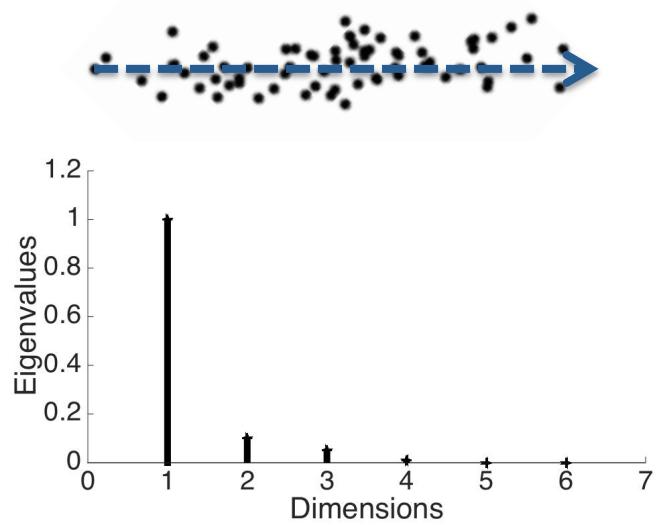
People usually use PCA to do 2 things:

1. Replot the data using  $Y$



2. Look at the spectrum (the eigenvalues)

Can tell you the ‘dimensionality of your data’

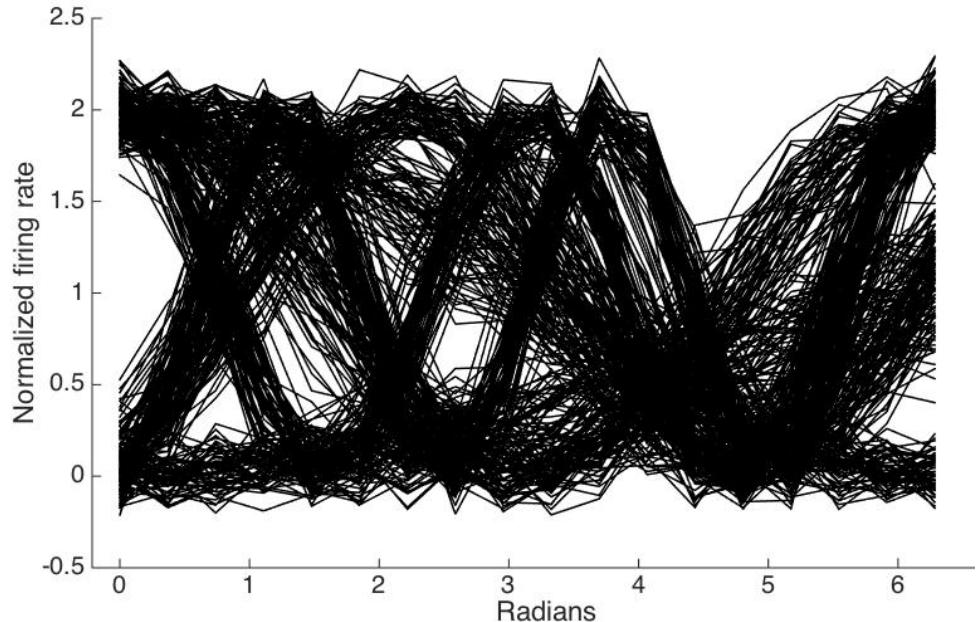


# PCA: how to do this in the real world

Step 0. Explicitly define what the data matrix X is.

Let's say we have a bunch (360) of head direction tuning curves.

We can tell there are a couple different types... but can we say this quantitatively?

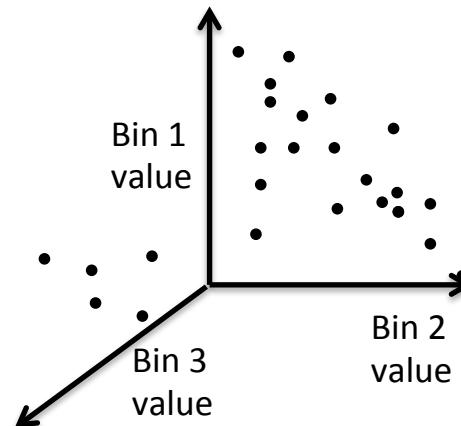


Bin each tuning curve into 18 bins

Each tuning curve is then a point in an 18-dimensional space

$$X = [TC1 \ TC2 \ \dots \ TC360]$$

X is a  $18 \times 360$  matrix



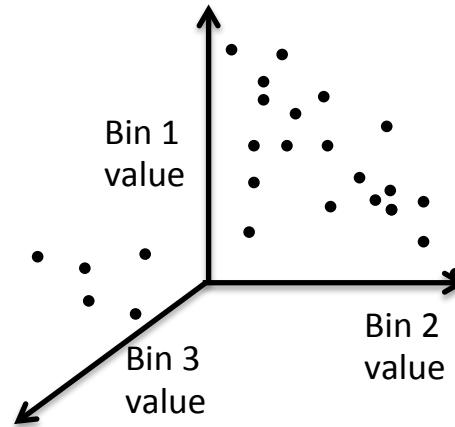
# PCA: how to do this in the real world

## Step 1. Center the data

- Compute and subtract off the mean
- Compute and divide by standard deviation

MATLAB:

```
[~,N] = size(X);  
mn = mean(X,2);  
X= X-mn*ones(1,N);  
stdX = std(X,[],2);  
X= X./(stdX*ones(1,N));
```



Otherwise, the first component will capture how far away from the origin the data is

## Step 2. Compute covariance matrix

MATLAB:

```
covariance = 1/(N-1)*X*X'; (or: covariance = cov(X'));
```

## Step 3. Compute and sort eigenvectors and eigenvalues of covariance matrix

MATLAB:

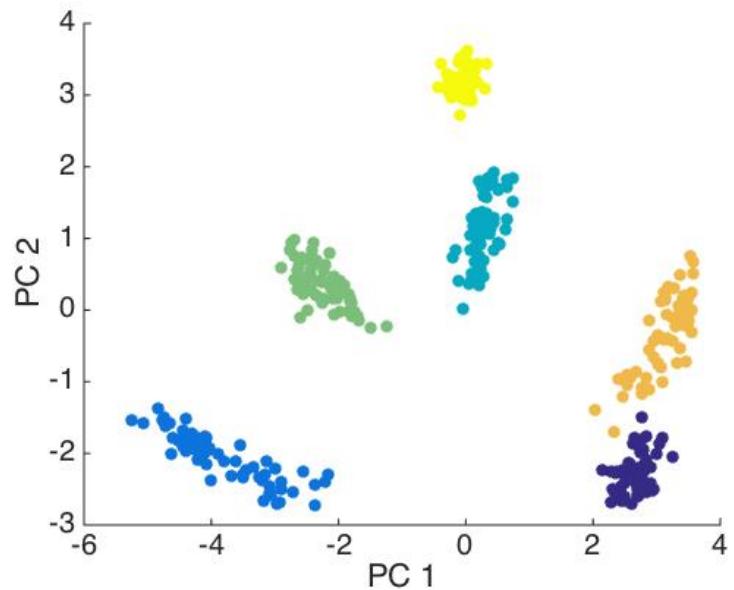
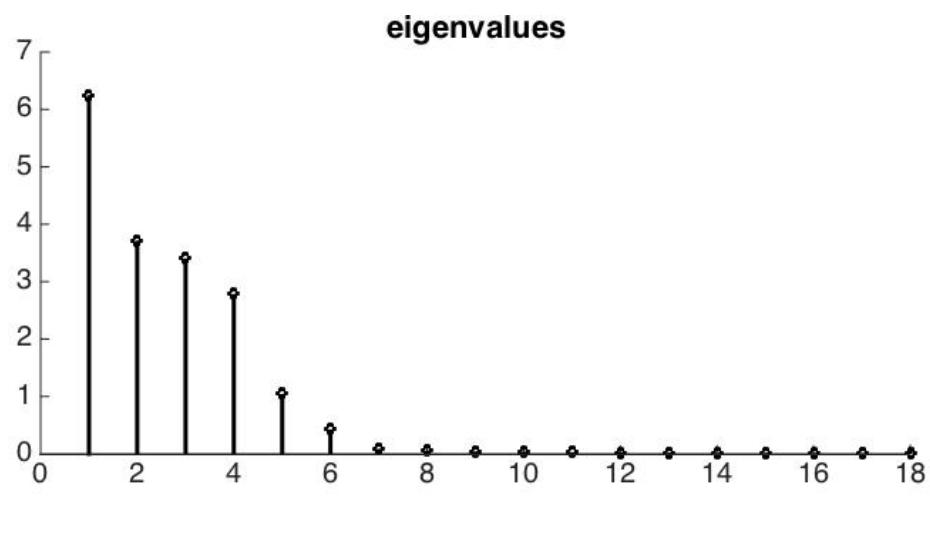
```
[PC,V]= eig(covariance);  
V = diag(V);  
[~,ind] = sort(V,'descend');  
PC = PC(:,ind); V = V(ind);
```

# PCA: how to do this in the real world

Step 4. Transform the data and see what you can see!

MATLAB:

```
Y = PC' * X;
```



7 types of cells

# When PCA fails

The linearity assumption is not a good one

- One fix: kernel PCA, where a nonlinear transformation is applied first

Finding the most decorrelated components isn't what you want

- ICA: independent component analysis: finds the statistically independent components

Interpreting dimensionality in the face of noise can be difficult

- Noise will increase the number of 'dimensions'
- Can be difficult to draw the line between signal and noise

# SVD

Recall that for some matrices (symmetric ones), we can decompose them into the following product:

$$X = VDV^{-1}$$

Where  $V$  is a matrix of eigenvectors and  $D$  is a matrix of eigenvalues along the diagonal

The SVD (singular value decomposition) is similar, but more general – you can do it for all matrices!

$$X = U\Sigma V$$

$U$  is a matrix of eigenvectors for  $XX'$  (the principal components!),  $V$  is a matrix of eigenvectors for  $X'X$ , and  $\Sigma$  is a matrix with singular values along the diagonal. Singular values are the square roots of  $X'X$

→ Columns of  $U$  are ‘left singular vectors’, row of  $V$  are ‘right singular vectors’

# SVD and applications

One great thing about the SVD: it decomposes a matrix into a sum of outer products

$$X = U\Sigma V$$

$$= \begin{pmatrix} | & | & & | \\ \vec{u}^{(1)} & \vec{u}^{(2)} & \dots & \vec{u}^{(N)} \\ | & | & & | \end{pmatrix} \begin{pmatrix} \lambda_1 & 0 & \dots & 0 & 0 & \dots & 0 \\ 0 & \lambda_2 & \dots & 0 & 0 & \dots & 0 \\ \vdots & \vdots & \ddots & \vdots & 0 & \dots & 0 \\ 0 & 0 & \dots & \lambda_N & 0 & \dots & 0 \end{pmatrix} \begin{pmatrix} - & \vec{v}^{(1)} & - \\ - & \vec{v}^{(2)} & - \\ \vdots & \vdots & \vdots \\ - & \vec{v}^{(N)} & - \\ other & & \end{pmatrix}$$

$$= \begin{pmatrix} | & | & & | \\ \vec{u}^{(1)} & \vec{u}^{(2)} & \dots & \vec{u}^{(N)} \\ | & | & & | \end{pmatrix} \begin{pmatrix} - & \lambda_1 \vec{v}^{(1)} & - \\ - & \lambda_2 \vec{v}^{(2)} & - \\ \vdots & \vdots & \vdots \\ - & \lambda_N \vec{v}^{(N)} & - \end{pmatrix}$$

$$= \lambda_1 \begin{pmatrix} | \\ \vec{u}^{(1)} \\ | \end{pmatrix} (- \vec{v}^{(1)} -) + \lambda_2 \begin{pmatrix} | \\ \vec{u}^{(2)} \\ | \end{pmatrix} (- \vec{v}^{(2)} -) + \dots + \lambda_N \begin{pmatrix} | \\ \vec{u}^{(N)} \\ | \end{pmatrix} (- \vec{v}^{(N)} -)$$

# SVD and applications

$$= \lambda_1 \begin{pmatrix} | \\ \vec{u}^{(1)} \\ | \end{pmatrix} (- \vec{v}^{(1)} -) + \lambda_2 \begin{pmatrix} | \\ \vec{u}^{(2)} \\ | \end{pmatrix} (- \vec{v}^{(2)} -) + \dots + \lambda_N \begin{pmatrix} | \\ \vec{u}^{(N)} \\ | \end{pmatrix} (- \vec{v}^{(N)} -)$$



Each term is a matrix

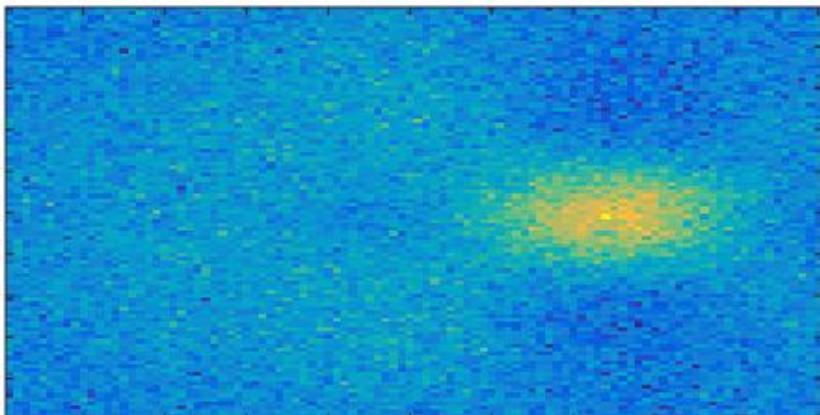
Each matrix is ordered in terms of ‘importance’

One application: compute low-rank approximations of matrices

MATLAB:

```
[U,S,V] = svd(X,0); % take an svd  
X_0 = U(:,1:p)*S(1:p,1:p)*V(:,1:p)';
```

Noisy receptive field:



Low-rank approximation:

